

Cohost

Aria - your podcast co-host

Version 1.0

Live tester edition - Hans (NVIDIA build)

Welcome, Hans. This is the first public build of Cohost. Aria is a podcast co-host that listens during your show and chimes in when you invite her. She is meant to be a second voice - she will not interrupt unless asked.

You are the live NVIDIA tester for this version. If anything does not work the way you expect, anything is missing, or anything feels wrong - tell cypherstone. He will track every request and fix them. There is a feedback form on the last page of this guide, and a one-click 'Remote help' button in the app that lets him connect to your machine to diagnose problems.

How to reach cypherstone

Email: ralstonjeff1978@gmail.com

In the app: click 'Remote help' in the sidebar - read him the URL and code.

Install

Prerequisites (one-time, ~5 min)

Before running the installer, get these three things on your PC:

- Python 3.12 from <https://www.python.org/downloads/release/python-3120/> - IMPORTANT: check 'Add python.exe to PATH' during install.
- Ollama for Windows from <https://ollama.com/download/windows>. After install, open PowerShell and run 'ollama signin' once to connect to Ollama Cloud. The \$20/month plan unlocks the bigger brains; free tier works with limits.
- NVIDIA driver - up-to-date if you can play modern games. If Aria runs slow later, update from <https://www.nvidia.com/Download/index.aspx>.

Install (one-time, ~15 min)

Once prerequisites are done:

- Make sure the Cohost folder has at least 15 GB free for downloads and model weights.
- Double-click INSTALL.bat. A window opens and shows progress. Do not close it.
- When you see 'Install complete' you are done with this step.
- If anything failed, email install\install_log.txt back to cypherstone.

Self-test (one-time, ~2 min)

This is the trust check before you record a real episode.

- Double-click SELFTEST.bat.
- It boots Aria, generates a test sentence, and lists your mics.
- Email install\selftest_report.txt back to cypherstone.
- He will green-light it before your first episode.

Running Aria

Open the app

Double-click CohostWeb.bat. Your browser opens to `http://localhost:6500` (first free port in 6500-6510) with the visual interface. This is what you will use 99% of the time.

There is also Cohost.bat for a terminal-only version. Same engine, no UI, slash commands instead of buttons.

How Aria behaves

Aria is a co-host, not the host. She listens passively until invited.

Wake word

Say 'Aria' or 'Hey Aria' during a recording. She opens a 90-second active window where she chimes in freely.

Hotkey

Press Ctrl+Alt+A from anywhere - works even when the browser is not in focus. Same effect as the wake word.

Cooldown

When the 90-second window closes, Aria enters a 180-second cooldown. During cooldown only a wake word reopens the window. This stops her from talking over you.

Free chime mode

Sidebar toggle. When ON, Aria replies to every turn. Useful for solo recordings where you want a constant conversation partner.

Aria, chime in (button)

Forces an immediate response from her without waiting for a wake word.

Sidebar controls

- Brain dropdown - top of sidebar. Grouped: 'Local (installed)' shows whatever you pulled with 'ollama pull'; 'Cloud' shows the curated list. Switch on the fly.
- Voice output toggle - mute Aria's voice but keep her text on screen.
- Mic selector - pick which microphone she listens on.
- Dream / Reflect - end-of-session reflection. Aria thinks about what went well and what she could do better. Improves her over time.
- Remote help - opens a 1-hour Cloudflare tunnel and shows you a URL plus a one-time code. Read both to cypherstone over the phone. Tunnel auto-closes after 1 hour or when you click Stop remote.
- End Episode - saves the session and closes Aria.

Reference

CLI slash commands (Cohost.bat only)

/brain	show brain menu, pick a different LLM
/brain <name>	switch brain directly
/voice	list voice options
/voice <name>	switch voice
/mic	list microphones
/mic <number>	switch to a specific mic
/whisper <size>	change Whisper STT model (tiny base small medium large-v3)
/guest <name>	set or change the guest
/bio <text>	set guest bio
/quote	save last guest line as a memorable quote
/followup <text>	flag a thread for follow-up
/open	generate an episode opening
/question	generate an interview question
/status	show system status
/memory	show memory stats
/learning	show learning report
/guests	list all past guests
/episodes	list recent episodes
/end	end and save current episode
/help	full command list
/quit	save and exit

Troubleshooting

Symptom	Fix
'Python 3.12 not detected'	Install Python 3.12, check 'Add to PATH' during install, re-run INSTALL.bat.
'Ollama not found'	Install from ollama.com/download/windows , then run 'ollama signin'.
'torch.cuda.is_available() is False' but you have NVIDIA	Update NVIDIA driver, re-run INSTALL.bat.
Aria takes 30+ seconds to respond	TTS fell back to CPU. Check selftest_report.txt - should say device = cuda. If not, update NVIDIA driver.
Mic not picking up	Web UI: change mic in the sidebar dropdown. CLI: /mic then /mic <number>.
Aria never responds	By default she is passive. Say 'Aria' or press Ctrl+Alt+A to invite her. Or toggle 'Free chime mode' in the sidebar.
Remote help button errors	Install cloudflared: 'winget install Cloudflare.cloudflared' and try again.
Anything else	Click 'Remote help' in the sidebar and send cypherstone the URL plus code.

This is v1.0 - your feedback shapes v1.1

You are the live NVIDIA tester for this build. Treat anything that surprises you as a bug report - cypherstone wants to hear it. Nothing is too small.

How to ask for changes

There is no formal process. Email, text, call - whatever works. Use the template below if it helps, or just describe what you saw and what you wanted to happen. Every request gets tracked and triaged.

Feedback template (copy / paste into an email)

Date:

Tester: Hans

Build: Cohost v1.0

Things that did not work

- [what happened, what you expected]
- [steps to reproduce if you can]

Things I want added

- [feature, why you want it]

Things I want changed or removed

- [what, why]

How Aria sounded / felt during the test

- [too slow, too fast, awkward pauses, talked over you, etc.]

Selftest result

- [device = cuda OR cpu, first_generate_secs, second_generate_secs]

Anything else

Examples of useful feedback

- 'When I say her name mid-sentence she sometimes responds before I finish. Can the wake-word be a full phrase like Aria-comma?'
- '180-second cooldown feels too long for a fast-paced segment. Make it configurable.'
- 'Brain dropdown is too small to read - bigger font please.'
- 'Want a button to save the last 30 seconds as a quote.'

Known limitations in v1.0

- Cloudflared (Remote help) is optional - installer tries winget but may need a manual install on

some systems.

- Self-test does not yet verify the new wake-word and remote-help code paths - selftest_report.txt may pass even if those pieces are subtly broken. Report any weirdness anyway.
- Hotkey (Ctrl+Alt+A) registration is best-effort - on some Windows configurations it needs admin privileges. If the hotkey does not work, the wake word and the 'Aria, chime in' button still do.

Bottom line

If something does not work, ask. If something is missing, ask. If something is wrong, ask. cypherstone will work on every issue you raise.